





KINESIS @TOTANGO

Ran Tavory

WHAT IS KINESIS

- Stream processing - as a service
 - Similar to Kafka
- Use cases:
 - Realtime data stream processing
 - Aggregations
 - Moving time-window, etc..





HIGH LEVEL

Overview




CONCEPTS

- Data Record
- Stream
- Partition Key
- Shard
- Sequence Number
- Worker



WRITING




WRITING

```
PutRecordResult putRecordResult = client.putRecord(putRecordRequest)  
// OR:  
client.putRecordAsync(putRecordRequest, asyncHandler)  
  
// In general:  
PutRecord(Data, PartitionKey, StreamName)
```



READING



READING

- Ordered
- Sequence Numbers
- Checkpoints
- “At least once” semantics
- Replay-ability (AT or AFTER _SEQUENCE_NUMBER)
- Iterator types:
 - AT_SEQUENCE_NUMBER, AFTER_SEQUENCE_NUMBER, TRIM_HORIZON (first avail), LATEST (most fresh)



KCL (READING)

- Kinesis Client Library
 - Java: <https://github.com/awslabs/amazon-kinesis-client>
 - (there are also for node, python, ruby)



KCL (READING)

- **Connects** to the stream
- **Enumerates** the shards
- **Coordinates** shard associations with other workers (if any)
- Instantiates a **record processor** for every shard it manages
- **Pulls data** records from the stream
- **Checkpoints** processed records
- **Balances** shard-worker associations when the worker instance count changes
- **Balances** shard-worker associations when shards are split or merged



KCL (JAVA)

```
// Processor (worker)
public interface IRecordProcessor {
    void initialize(String shardId);
    void processRecords(List<Record> records, IRecordProcessorCheckpointer checkpointer);
    void shutdown(IRecordProcessorCheckpointer checkpointer, ShutdownReason reason);
}

// Factory
public interface IRecordProcessorFactory {
    IRecordProcessor createProcessor();
}
```



RE-SHARDING

- So what's the story with re-sharding?
 - Choosing the partition key
 - Shard limits
 - 1Mb/S ingest, 2Mb/S egress, 1K inserts/S
 - Resharding - It is painful :-(



RE-SHARDING (EXAMPLE)

```
$ aws kinesis describe-stream --stream-name  
gateway-received
```



```
$ aws kinesis split-shard --stream-name  
gateway-received --shard-to-split  
shardId-000000000017 --new-starting-hash-key  
255211775190703847597530955573826158591
```



KINESIS V/S SQS

- Speed
- Data item size
- Ordered messages
- Replay-ability
- Sharding (persistent routing)
- Kinesis workers - v/s SQS consumers



KINESIS CONNECTORS

- <https://github.com/awslabs/amazon-kinesis-connectors>
- Amazon DynamoDB
- Amazon Redshift
- Amazon S3
- Elasticsearch



PRICING

Pricing Dimension	Value
Hourly Shard Rate	\$0.015
Per 1,000,000 PUT transactions:	\$0.028

- Shard / month => ||\$




TOTANGO CLASSES

```
/**  
 * A high-level client for Kinesis  
 * @author ran  
 */  
public class KinesisClient {  
  
    public void connect() {  
  
        public List<String> listStreams() {  
  
            public boolean isStreamExists(final String streamName)  
  
            public String describeStream(final String streamName)  
  
            public List<Shard> describeStreamShards(...)  
  
            public List<Shard> describeStreamLeafShards(...)  
  
            public void createStream(...)  
  
            public String putRecord(...)  
  
            public void putRecordAsync(...)  
    }
```



KINESIS @ TOTANGO



REFERENCES

- <http://aws.amazon.com/kinesis/>
- <http://docs.aws.amazon.com/kinesis/latest/dev/developing-consumer-apps-with-kcl.html>
- <https://github.com/awslabs/amazon-kinesis-client>
- <https://github.com/awslabs/amazon-kinesis-connectors>
- This presentation: <https://speakerdeck.com/rantav/kinesis>

