



Lambda Stories from the Trenches

What worked and what didn't work for us using AWS Lambda

@rantav
Yodas
(now AppsFlyer)

whoami

- Yodas cofounder and Janitor
- Now at AppsFlyer
- Reversim Podcast
- Reversim Summit (Oct 8-9)
- ... been coding Lambda since early 2016

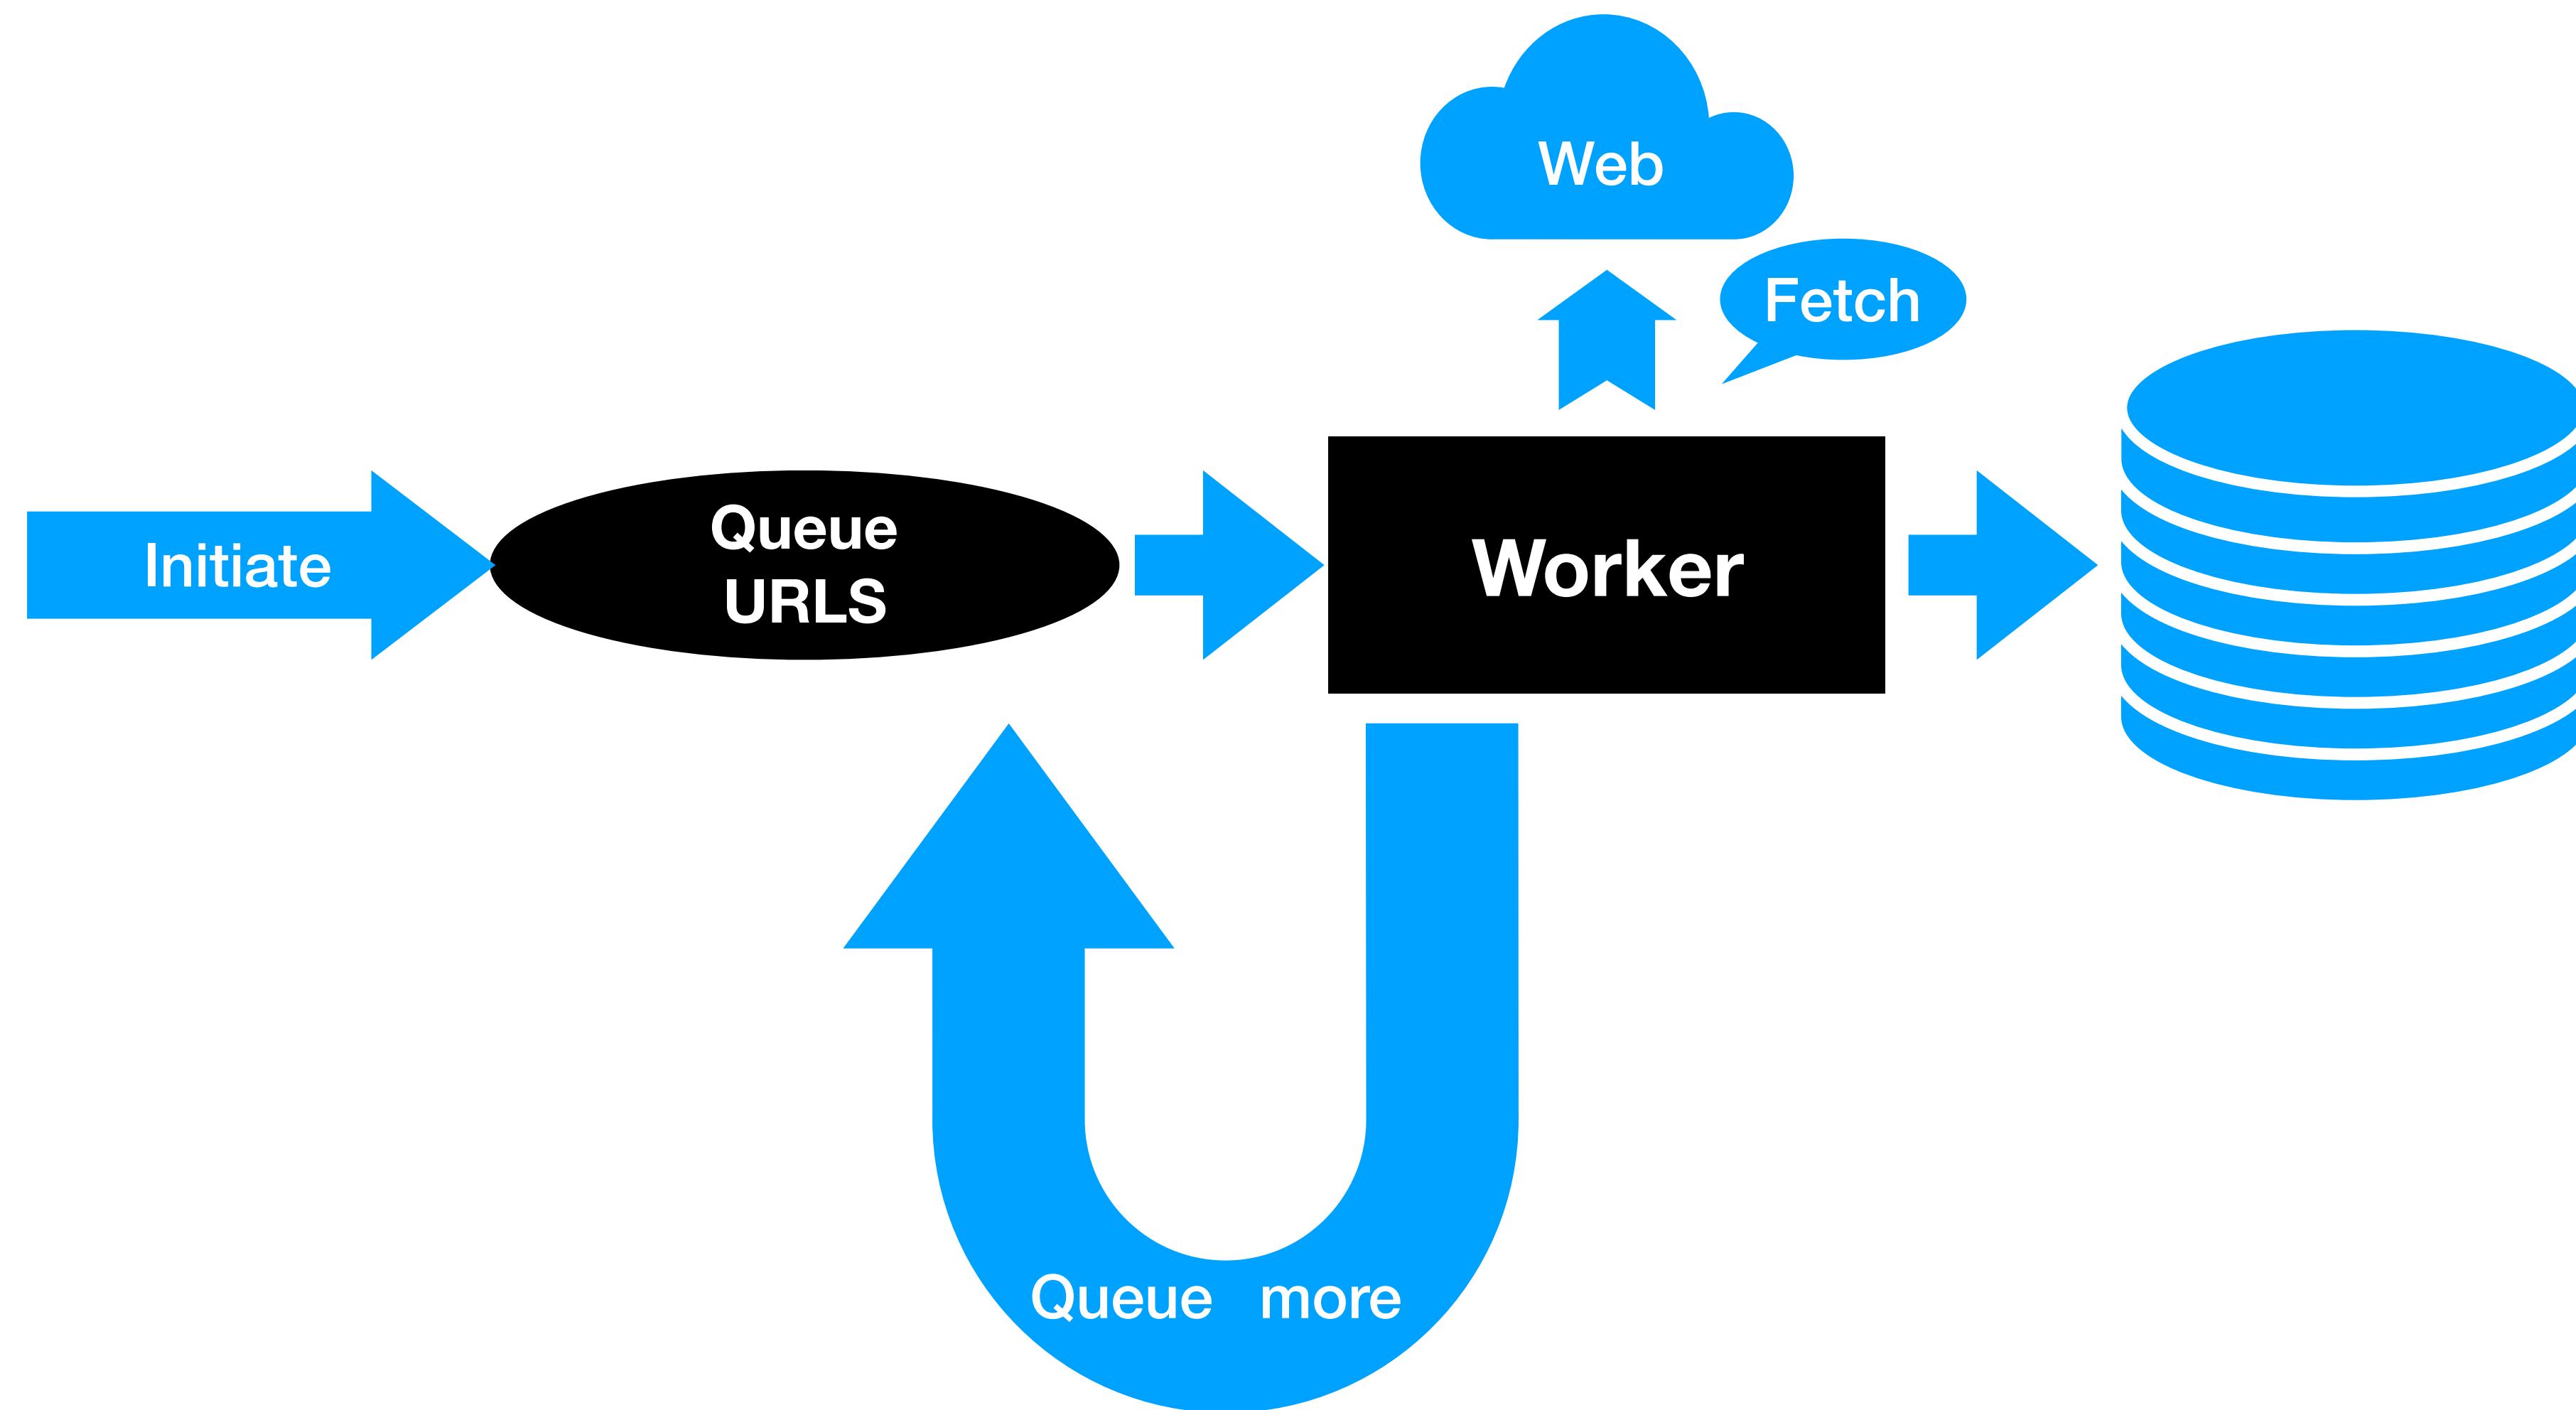


A tale of two Lambdas

- The Happy Lambda
- The Sad Lambda

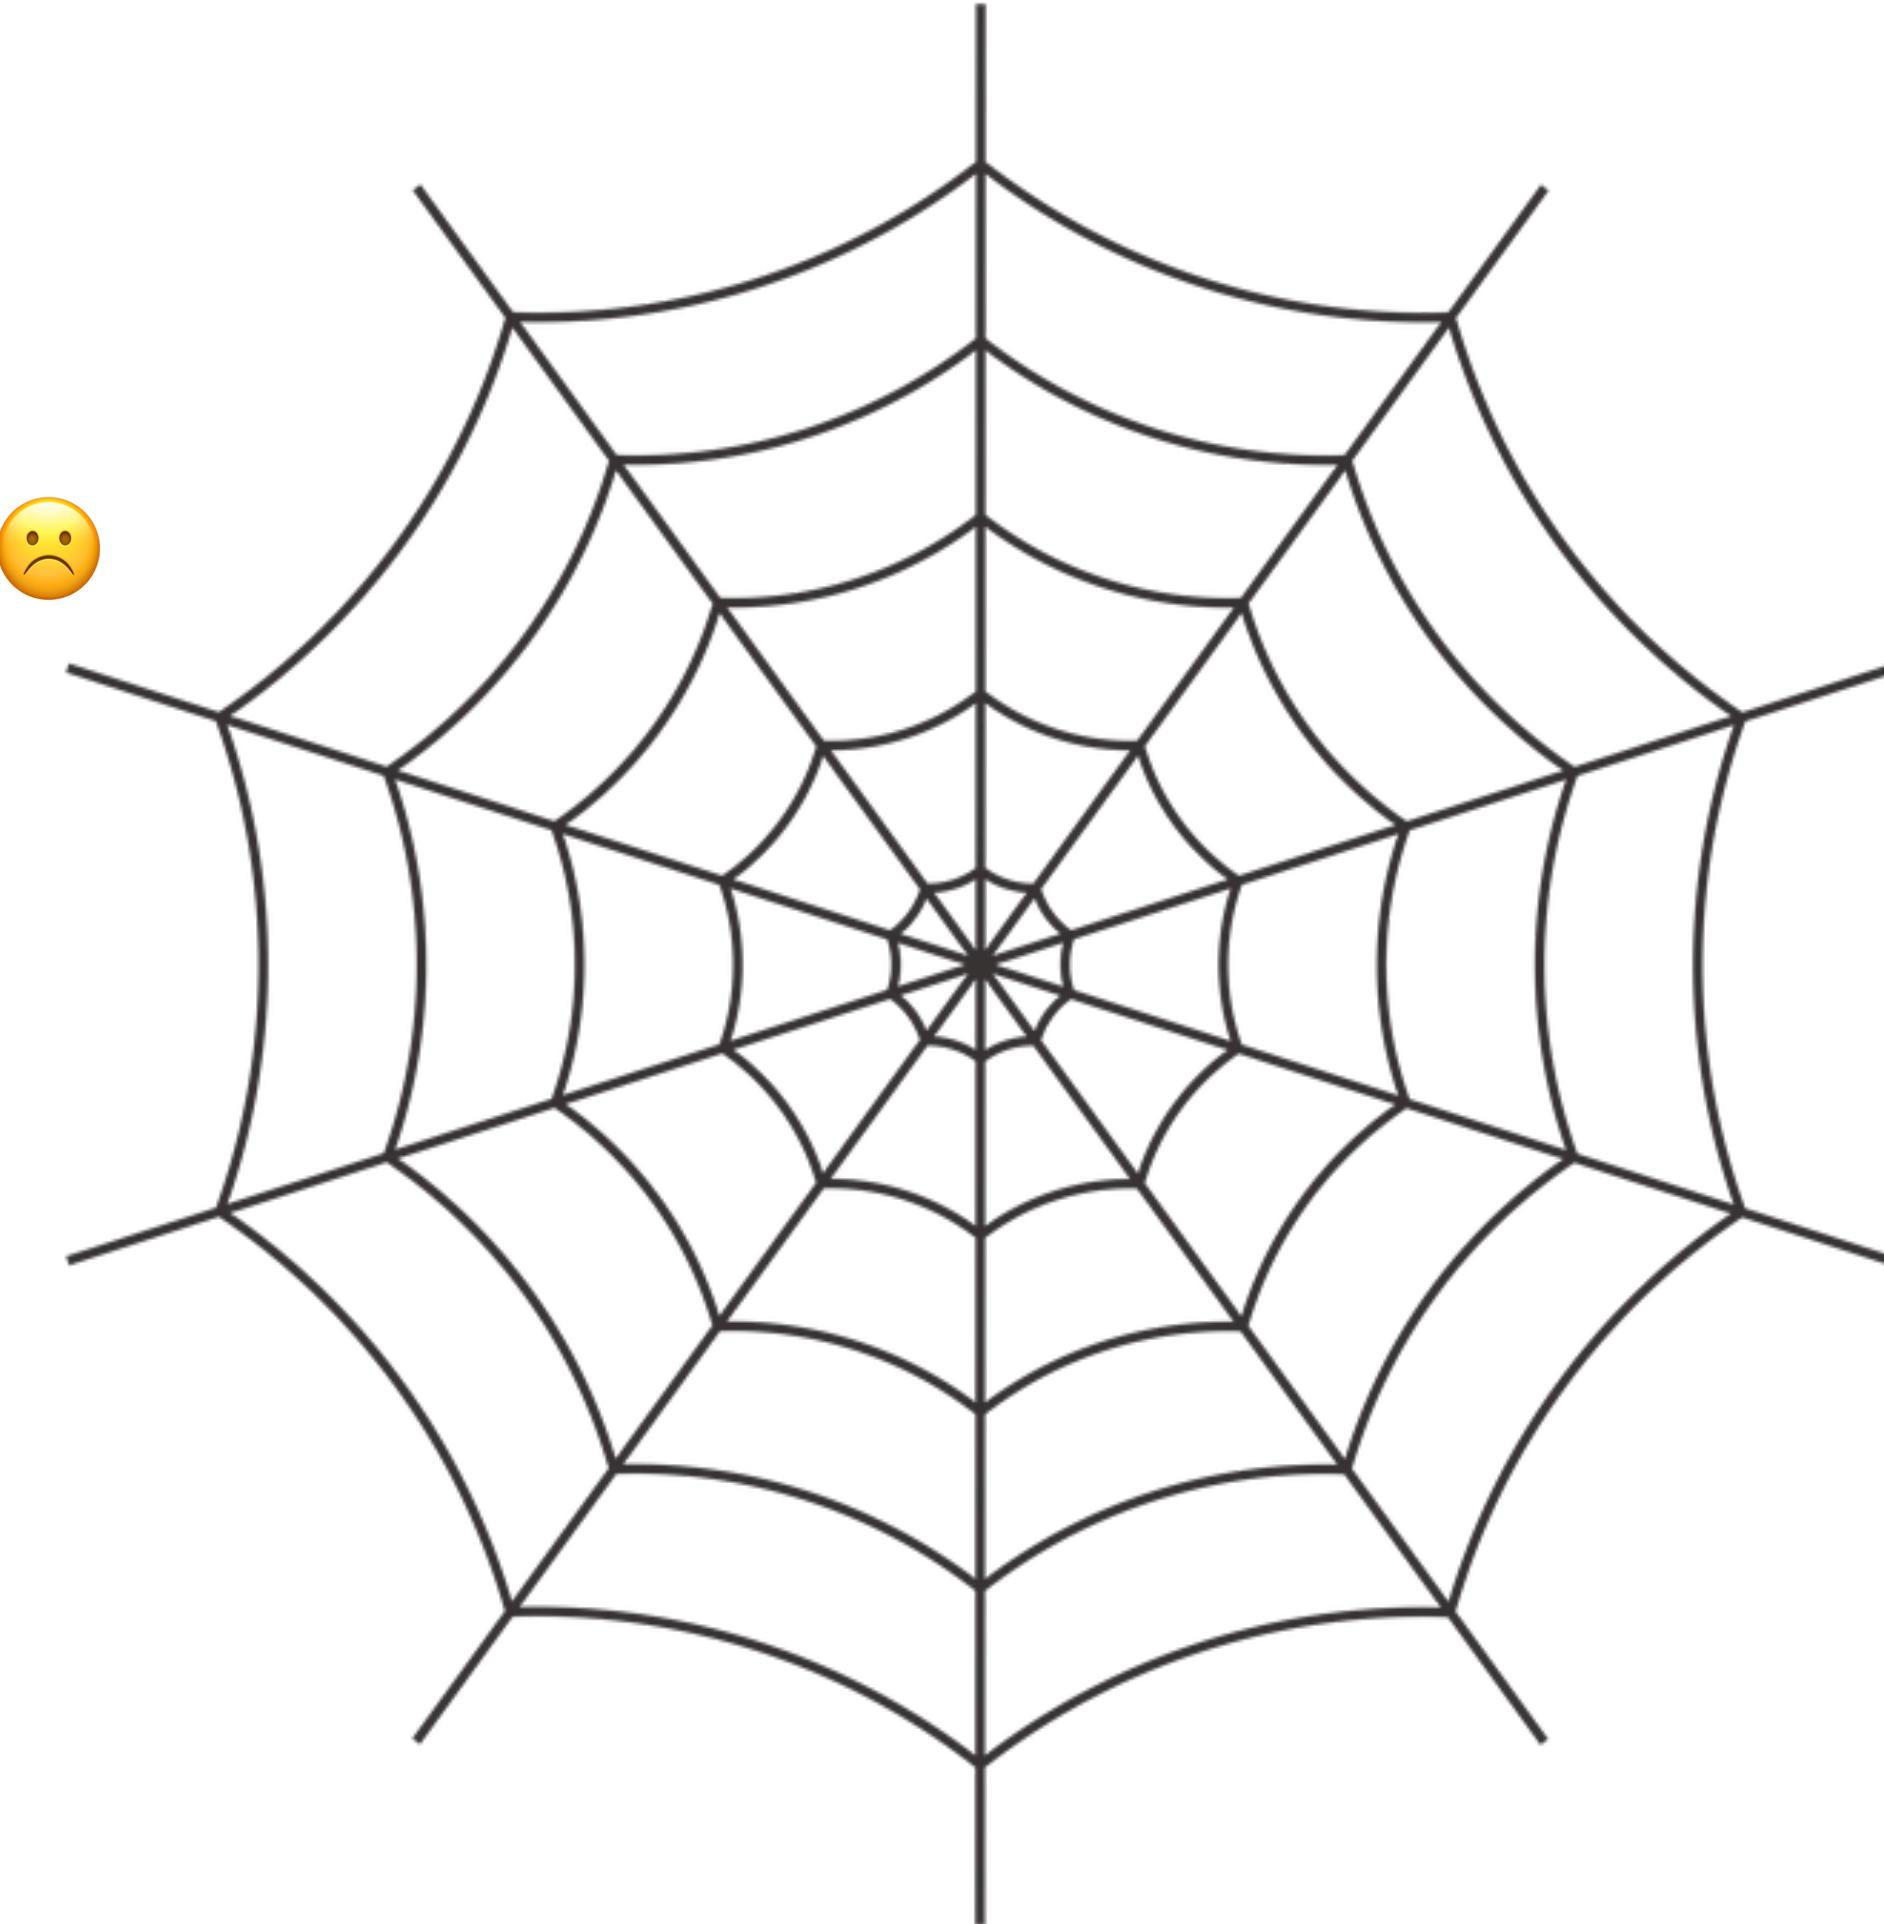


How a web crawler works (simplified)



Web Crawler Properties

- Typical page: seconds → sometimes minutes 😕
- Failures / Retries ↗
- Easily distributed 👍
- Throttle 🤐



Web Crawler on AWS Lambda

- Easy 🤘
- Function: Fetch
 - Input: URL
- But wait, no SQS support!
- Go + Apex



Shortcomings

- Async invocation 🤢
- Throttling 🤢
- Concurrency 🤢



Crawler on Lambda Take 2

- Concurrency:
 - Micro Batching
 - But rough edges



Take 2

Crawler on Lambda: Challenges

- Cost
- Throttling
- SQS



Crawler on Lambda: Conclusion

FAIL



Solution: Elastic Beanstalk workers

- Worker mode works with SQS
- Spots
- Scale and Concurrent
- Capacity and Throttling support.



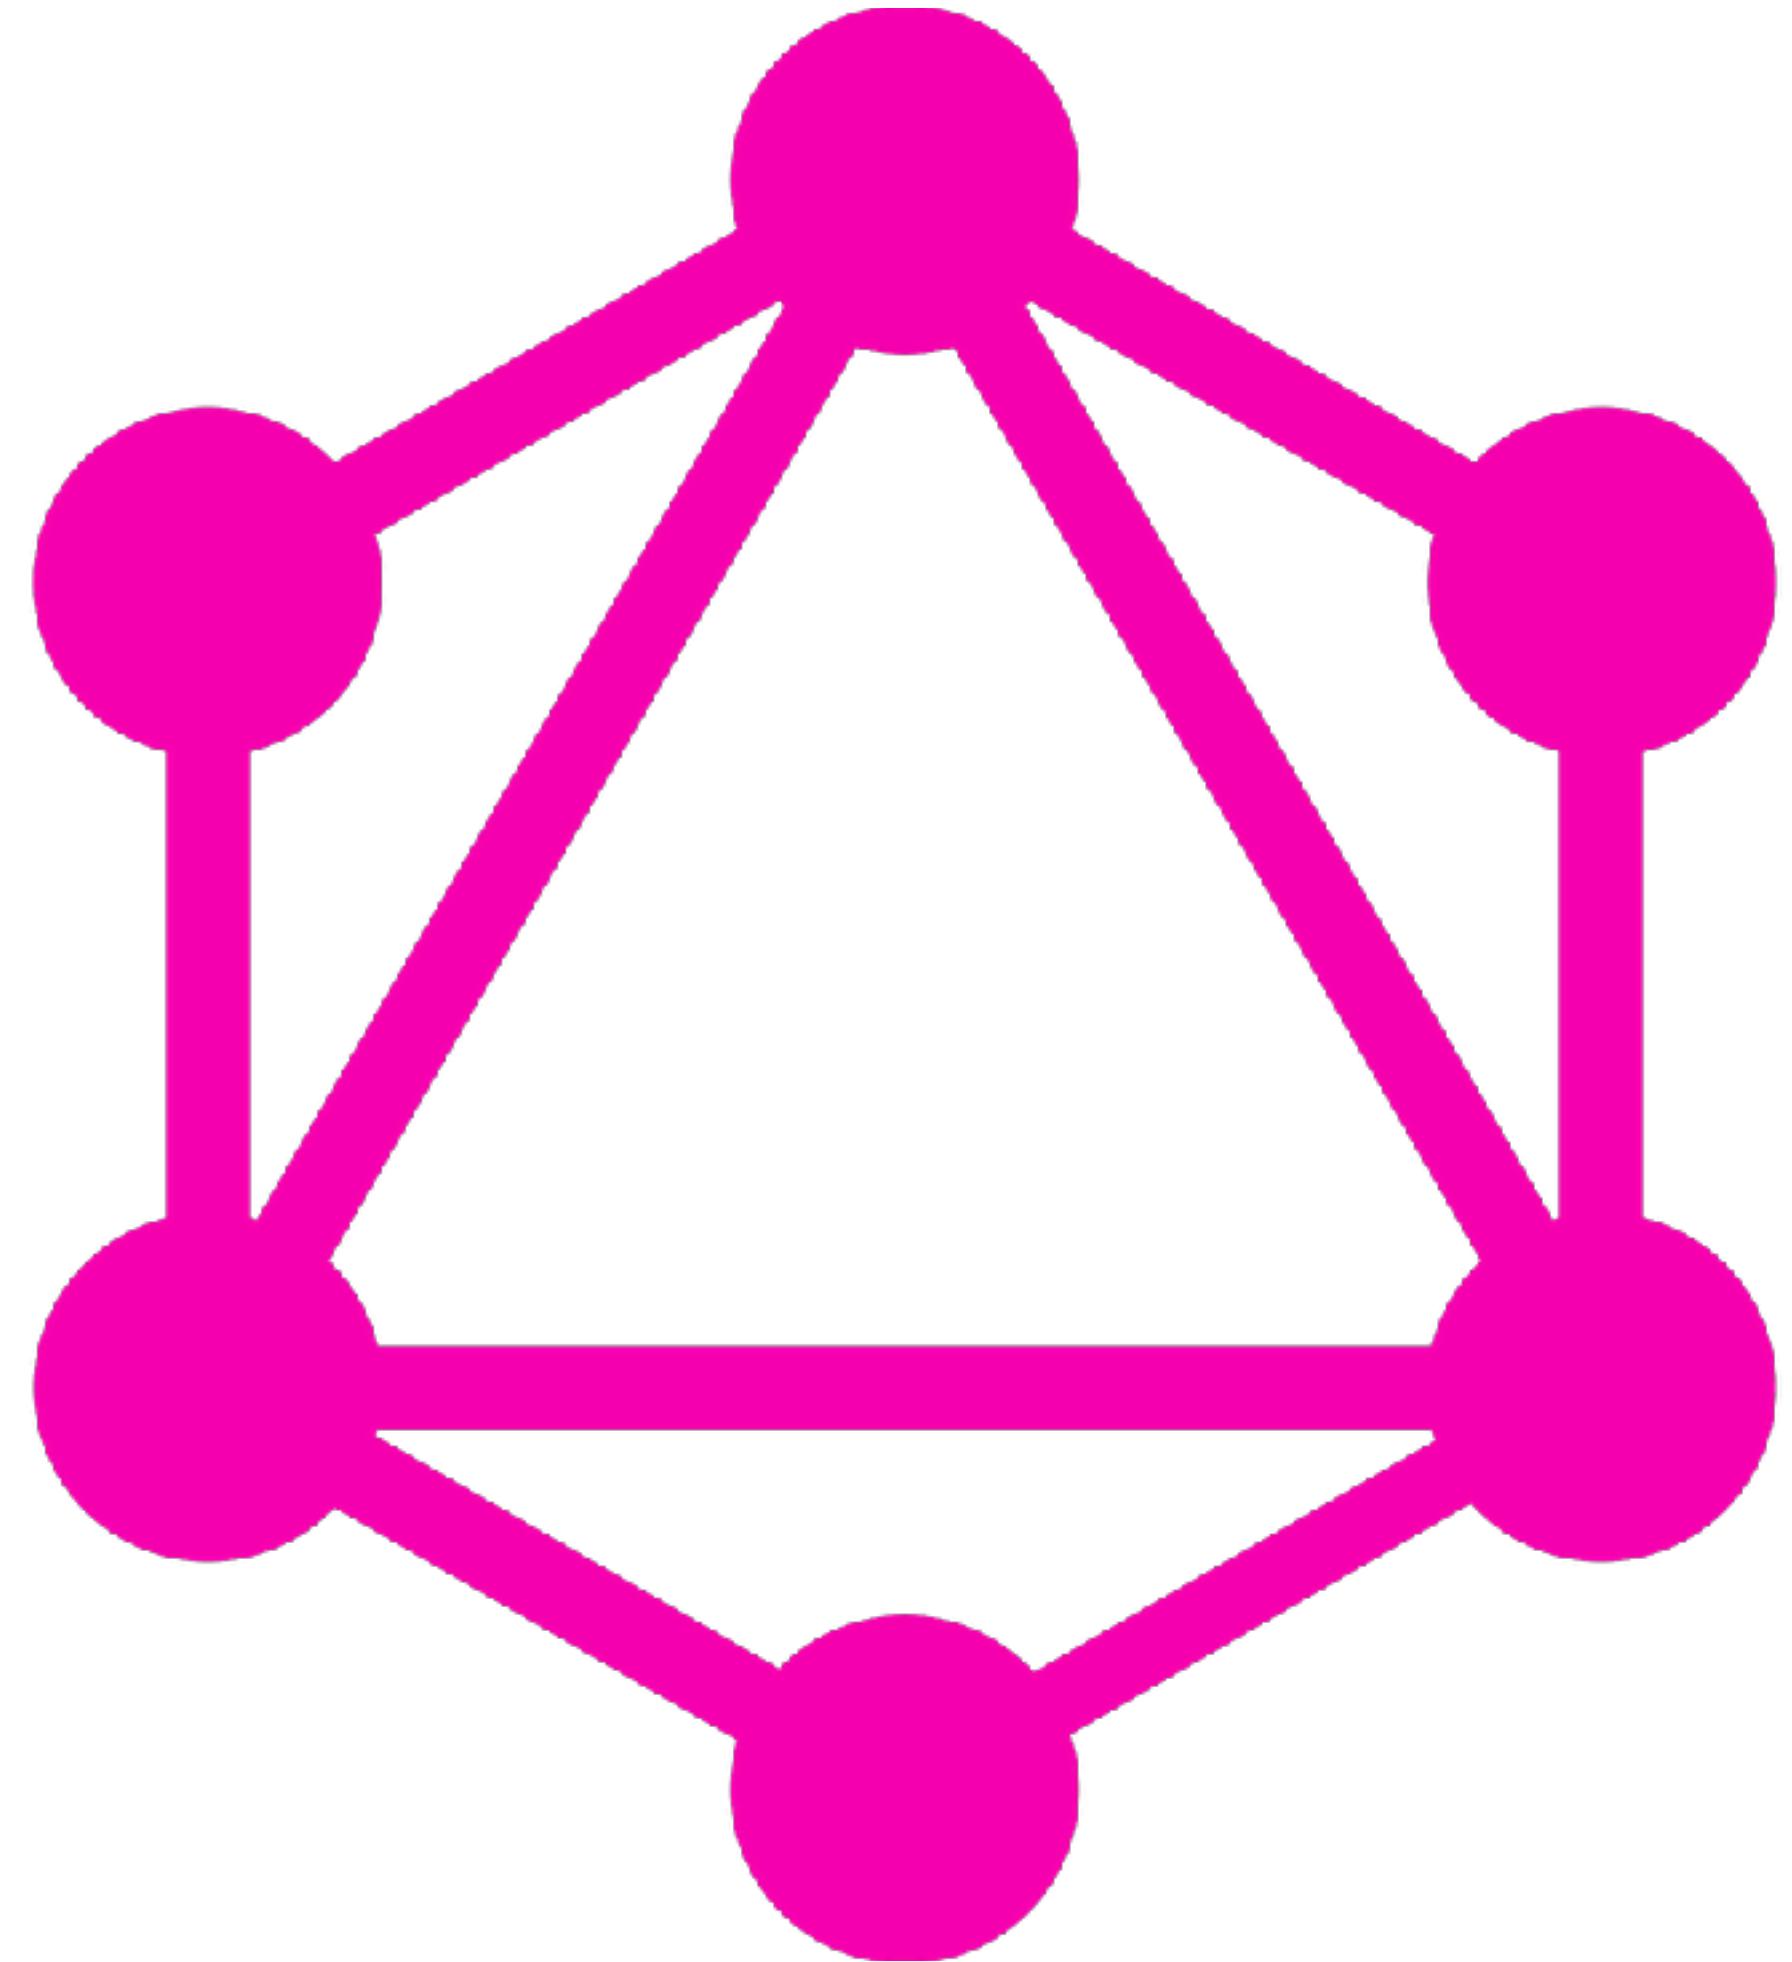
API on Lambda

- The stack:
 - API GW
 - Go + Apex
 - Terraform
 - GraphQL



Why GraphQL?

- Clients love it
- Easy to maintain a monolith
 - Wait, what?



Observation: Ops @ Lambda

- NoOps?
- No, Ops



API on Lambda: Challenges

- Latency
 - Cold start
 - API GW
 - Conn pooling
- Visibility
 - Logging
 - Monitoring



Latency Mitigation

- Delayed/background jobs
- Conn pooling when possible
- No premature initialization



Conclusion

- Lambda sounds good, in theory
 - Reality: practical limitations
- GraphQL is a nice fit for API atop FaaS
- API: If you can afford the latency, it's worth it

